

# **Siddes Observability Pack**

Production debugging that is fast, privacy-first, and beginner-friendly.

Generated on 2026-01-25

# Table of Contents

Table of Contents	2
Siddes Observability Spec v1 (Privacy-first, Launch-grade)	3
Siddes Observability - ClickOps Implementation (Datadog)	7
Siddes - Datadog Dashboards & Monitors Pack (Copy/Paste)	12
Siddes Observability Fire Drill (30 minutes, beginner-proof)	19
Siddes Observability - Vercel/Next Pack (Inbox Stub + /api/health)	22
Siddes Support Code Workflow (No Code Changes)	25
Siddes Support Response Macros (Copy/Paste)	28
Siddes Incident Note Template (Copy/Paste)	30
Siddes Oncall Home (Single Pane)	32
Siddes Oncall Quickstart (Beginner)	35
Siddes Observability Master Checklist (Rebuild From Scratch)	36
Observability (Production Debugging, Privacy-First)	40
Siddes Burn-Rate Alerts (Smarter paging, less noise)	44
Siddes - Fix Playbook Map (Signature → Evidence → Repo Knobs → Verify)	48
Siddes Observability Incident Drill Pack (Beginner, 45 minutes)	57
Siddes Known Failure Signatures (Datadog “What am I seeing?” Guide)	61
Siddes Launch Readiness - Observability Gate (Run Before Every Prod Deploy)	66
Siddes Observability Privacy Audit Checklist (Monthly)	68
Siddes Observability Retention & Access Policy (v0)	72
Siddes Observability SLOs v0 (No Code Changes)	74
Siddes Observability Pack Index (Read Order)	77
Siddes Oncall Observability Playbook (Beginner-friendly)	79

# Siddes Observability Spec v1 (Privacy-first, Launch-grade)

File: docs/OBSERVABILITY\_SPEC\_V1.md (source: sd\_600\_observability\_docs\_datadog\_apply\_helper\_20260125\_053723.sh)

This spec is designed to make production debugging fast **without** logging user content or identities.

It assumes Siddes' current topology:

- **Frontend:** Vercel (Next.js)
- **Backend:** DigitalOcean App Platform (Django)
- **Media:** Cloudflare R2 + optional Cloudflare Worker on /m/\*

Siddes already provides a strong baseline:

- Every /api/\* response includes X-Request-ID.
- Backend emits structured JSON-line request logs (event="api\_request") for /api/\*.
- Next /api proxy generates/forwards x-request-id and often injects requestId into JSON responses.

## 1) Goals (what we must answer in 5 minutes)

### Errors

- Are 5xx errors spiking? Which endpoint(s)? Which service (Vercel vs Django vs Cloudflare Worker)?

### Latency

- Is p95 latency spiking? Which endpoint(s)? Is it regional or global?

### Auth failures

- Are 401/403 spikes happening (especially for write endpoints)? Did cookies/CSRF break?

### Media failures

- Are sign-upload or commit failing? Are /m/\* reads failing at edge?

### Inbox failures

- Is send latency high? Is send failing? Is the failure in Next (Vercel) or Django?

## 2) Logging Standard (privacy-safe)

Siddes backend currently emits:

- event, request\_id, viewer, method, path, status, latency\_ms, side, (and sometimes query)

### Required stored fields (canonical)

Store these fields for every request log event:

- ts (timestamp)
- event (normalized to api.request)
- service (siddes-backend, siddes-next, siddes-edge)
- env (prod|staging)
- release (deploy SHA/version)
- request\_id
- method
- route (template, not raw IDs; e.g. /api/post/:post\_id/like)
- status, status\_class (2xx|3xx|4xx|5xx)
- latency\_ms
- side (public|friends|close|work|unknown)
- client\_version (frontend build/version)

### PII rules: NEVER log

- Passwords, OTPs, session cookies, auth/refresh tokens, Authorization header
- Email, phone, username, real name
- Post bodies, comments, inbox/DM text, bios, search queries
- Signed URLs, media keys, raw IP addresses

### Identity handling (strict)

- Do **not** index/store raw viewer values.
- Either **drop** viewer or convert to a non-reversible hash (preferred).

## 3) Correlation / Tracing Plan (MVP)

Siddes already has correlation through request\_id:

- Client → Next (/api proxy) → Django
- x-request-id forwarded
- Django logs request\_id and returns X-Request-ID

### Support Code rule

- For API/action failures: use request\_id

- For UI error boundaries: use `Next error.digest` (if/when you add client telemetry lane)

“Like didn’t work” workflow: 1) Grab the `requestId` in the failing JSON or the `X-Request-ID` response header 2) Search logs by `request_id` 3) Read `route, status, latency_ms` 4) If no backend log exists → it failed before reaching Django (Vercel runtime / browser issue)

## 4) Metrics + Dashboards (log-derived)

Build log-based metrics from request logs:

- **Count:** requests by route/status/method/side
- **Distribution:** `latency_ms` by route

Dashboards: 1) **Golden Signals:** RPS, 5xx rate, p95 latency, top failing routes, top slow routes 2) **Endpoint Explorer:** route table (rps, p95/p99, 4xx/5xx), status breakdown, sample logs 3) **Auth & Session:** 401/403 trends, write-method 401 outside `/api/auth/*`, csrf/me checks 4) **Media Pipeline:** sign-upload + commit success/fail + p95; `/m/*` edge status breakdown 5) **Inbox Reliability:** send success/fail + p95; split by Next vs Django where applicable

## 5) Alerts (page vs notify)

### Page

- Global 5xx ratio > 1% for 5 minutes (or “5xx count > threshold” if you’re early stage)
- Core routes (auth/me, auth/csrf, feed, post create/like, media commit, inbox send) failing > 2% for 3 minutes
- Core p95 latency > 1500ms for 10 minutes
- Media commit failures > 5% for 10 minutes
- Health endpoints failing consecutively

### Notify

- Write-method 401 spike outside `/api/auth/*` (CSRF/cookie drift)
- 429 spikes (rate limit)
- Region-specific spikes

Every alert payload must include:

- top affected routes
- example `request_ids` to search immediately

## 6) Privacy-safe telemetry controls

- **Retention:** raw logs  $\leq$  7 days (recommended)
- **Indexing:** index errors + slow requests; minimize indexing of 2xx
- **Sensitive Data Scanner / redaction:** hash viewer IDs; redact query strings
- **Access tiers:**
- Oncall/SRE: raw logs + request\_id search
- Engineers: dashboards; time-bound raw access
- Support: no raw logs; uses support codes

## 7) Implementation (recommended: Datadog)

See `docs/OBSERVABILITY_CLICKOPS_DATADOG.md` for a copy/paste runbook to:

- forward DigitalOcean App Platform logs to Datadog
- drain Vercel logs to Datadog
- Logpush Cloudflare `/m/*` traffic logs to Datadog
- configure Sensitive Data Scanner to hash/redact
- build dashboards + alerts

# Siddes Observability - ClickOps Implementation (Datadog)

File: docs/OBSERVABILITY\_CLICKOPS\_DATADOG.md (source:  
sd\_600\_observability\_docs\_datadog\_apply\_helper\_20260125\_053723.sh)

This is the “do exactly this” setup to centralize Siddes logs + build dashboards/alerts, **without changing Siddes code.**

You will connect:

- DigitalOcean App Platform → Datadog Logs
- Vercel Drains → Datadog Logs
- Cloudflare Logpush (/m/\*) → Datadog Logs

## Step 1 - Create Datadog keys (5 minutes)

1) In Datadog: **Organization Settings** → **API Keys**

- Create an API key named `siddes-ingest`

2) In Datadog: **Organization Settings** → **Application Keys**

- Create an App key named `siddes-admin` (needed for some config APIs; safe to keep private)

## Step 2 - Forward backend logs (DigitalOcean App Platform → Datadog)

DigitalOcean App Platform supports forwarding logs to Datadog.

1) DigitalOcean Control Panel → **Apps** → select your Siddes backend app 2) Go to **Settings** 3) Find **Log Forwarding** → click **Edit** 4) Choose **Datadog** 5) Paste your Datadog API key and select the correct Datadog site (US/EU) 6) Select the backend component(s) to forward 7) Save

### Definition of Done

- In Datadog Log Explorer, search: `@event:api_request`
- You should see logs with `request_id`, `path`, `status`, `latency_ms`

## Step 3 - Drain Vercel logs (Vercel → Datadog)

Vercel Drains can send logs to Datadog. Drains are available on Pro/Enterprise.

1) Vercel Dashboard → **Team Settings** → **Drains** 2) Click **Add Drain** 3) Choose **Logs** 4) Destination:

- Prefer: **Datadog integration** (Marketplace) if available
- Otherwise: **Custom HTTP endpoint** pointing to your Datadog Logs intake endpoint

5) Format: **NDJSON** (recommended) 6) Environment: **Production** 7) Save

### Definition of Done

- In Datadog Log Explorer, search for Vercel logs and confirm you see fields like:
- `requestId`, `statusCode`, `path`, `executionRegion`

## Step 4 - Logpush Cloudflare /m/\* edge logs (Cloudflare → Datadog)

Goal: see whether /m/\* is returning 200/302/401/403/5xx and where.

1) Cloudflare Dashboard → **Logs** → **Logpush** 2) Create a Logpush job (zone-level for your app domain) 3) Destination: **Datadog** 4) Endpoint: use Datadog HTTP intake endpoint (v1 is fine) 5) API Key: your Datadog API key 6) Dataset: `http_requests` 7) Filters:

- include only requests where **path begins with** /m/

8) Tags (as destination params):

- `ddsource=cloudflare`
- `service=siddes-edge`
- `ddtags=app:siddes,env:prod`

9) Submit

### Definition of Done

- In Datadog logs, search: `service:siddes-edge @http.request.uri.path:/m/*`
- You should see status codes and edge metadata.

## Step 5 - Enable Sensitive Data Scanner (hash viewer + redact query)

Siddes backend logs include `viewer` (e.g. `me_123`) and may include `query`. We will **hash** `viewer` and **redact** `query` before it becomes searchable/indexed.

In Datadog: 1) **Organization Settings** → **Sensitive Data Scanner** 2) Create a **Scanning Group for Logs**

- Filter: `@event:api_request`
- (Optional) restrict to service `siddes-backend`

3) Add Rule: **Hash viewer**

- Scope: scan attribute `viewer`
- Regex pattern: `^me_[0-9]+$`
- Action: **Hash**
- Tag: `siddes_sensitive:viewer`

4) Add Rule: **Redact query**

- Scope: scan attribute `query`
- Regex pattern: `.+`
- Action: **Redact** with replacement text [`redacted_query`]
- Tag: `siddes_sensitive:query`

5) Save + wait a couple minutes

#### Definition of Done

- New logs show `viewer` as a hashed token (or redacted), not raw `me_123`.
- `query` shows [`redacted_query`] (or is empty).

### Step 6 - Normalize routes (avoid cardinality explosion)

We want route templates like `/api/post/:post_id/like` instead of raw IDs.

In Datadog → **Logs** → **Configuration** → **Pipelines** Create pipelines that match paths and set a constant route attribute via **String Builder**.

Create these (minimum set):

#### A) Post routes

- Filter: `@event:api_request @path:/api/post/*/like`
- Set `route = /api/post/:post_id/like`
- Filter: `@event:api_request @path:/api/post/*/reply`
- Set `route = /api/post/:post_id/reply`
- Filter: `@event:api_request @path:/api/post/*/replies`
- Set `route = /api/post/:post_id/replies`
- Filter: `@event:api_request @path:/api/post/*/quote`
- Set `route = /api/post/:post_id/quote`
- Filter: `@event:api_request @path:/api/post/*`
- Set `route = /api/post/:post_id`

#### B) Inbox routes

- Filter: `@event:api_request @path:/api/inbox/thread/*`
- Set `route = /api/inbox/thread/:thread_id`

#### C) Media (static already)

- `/api/media/sign-upload`
- `/api/media/commit`

(These don't need templating; just set `route=%{path}` if you want.)

### Definition of Done

- In logs, you can facet/group by `route` without seeing thousands of unique values.

## Step 7 - Create log-based metrics

In Datadog Log Explorer:

- 1) Query: `@event:api_request`
- Click **Generate Metric**

- Metric name: `siddes.api.request.count`
- Type: count
- Tags: `route,method,status,side,service,env`

- 2) Query: `@event:api_request`

- Generate metric from attribute: `latency_ms`
- Metric name: `siddes.api.request.latency_ms`
- Type: distribution
- Tags: `route,method,side,service,env`

## Step 8 - Dashboards (MVP)

### Dashboard 1: Golden Signals

- RPS: `siddes.api.request.count` (timeseries)
- 5xx: query logs or metric filtered where status is 5xx
- p95 latency: `p95(siddes.api.request.latency_ms)` overall + by route
- Top failing routes: table grouped by `route`
- Top slow routes: table grouped by `route`

### Dashboard 2: Media Pipeline

- sign-upload success/fail rate
- commit success/fail rate
- p95 latency for sign/commit
- Edge /m/\* status breakdown (from Cloudflare logs)

### Dashboard 3: Auth & Session

- 401/403 on `/api/auth/*`
- Write-method 401 outside `/api/auth/*`

## Dashboard 4: Inbox Reliability

- send success/fail and p95 on /api/inbox/thread/:thread\_id

## Step 9 - Alerts (MVP)

### Paging

- Backend 5xx count > 10 in 5 minutes (group by route)
- Media commit failures > 5 in 10 minutes
- Health checks failing (if you add synthetic checks)

### Notify

- Write-method 401 spike outside auth (cookie/CSRF drift)
- p95 latency > 1500ms for 10 minutes on core routes

### Alert payload must include

- top routes
- sample request\_id values

# Siddes - Datadog Dashboards & Monitors Pack (Copy/Paste)

File: docs/OBSERVABILITY\_DATADOG\_DASHBOARDS\_ALERTS.md (source:  
sd\_601\_observability\_datadog\_dashboards\_alerts\_apply\_helper\_20260125\_054406.sh)

You said you finished the Part 4 checks (logs are flowing, @event:api\_request is visible, privacy rules are active). Now do **exactly** this to get production-debuggable dashboards + alerts.

This pack is designed to work **without changing Siddes code**.

## 0) One-time setup in Datadog (2 minutes)

### Create facets (if not already)

In **Logs → Explorer**, find a recent @event:api\_request log. For each attribute below, click it and choose **Create facet**:

- route (string) \*(from your Datadog pipeline templating)\*
- path (string)
- method (string)
- status (number)
- side (string)
- latency\_ms (number)
- request\_id (string)

Facets let you group dashboards by route/status and drill down fast.

## 1) Create 2 log-based metrics (skip if already created)

Go to **Logs → Explorer** and run this query:

### A) Requests count metric

Query:

```
@event:api_request
```

Click **Generate Metric**:

- Metric name: siddes.api.request.count
- Type: count
- Tags to keep (important): route,method,status,side
- (Optional if you have them): service,env,release

## B) Latency distribution metric

Query:

```
@event:api_request
```

Click **Generate Metric**:

- Metric name: siddes.api.request.latency\_ms
- Type: distribution
- Measure: attribute latency\_ms
- Tags to keep: route,method,side
- (Optional): service,env,release

Why distributions: you can do p95( . . . ) cleanly.

## 2) Dashboard Pack

Create a new dashboard for each section below.

### Dashboard 1 - Siddes: Golden Signals (Prod)

#### Widget 1: RPS (backend)

Type: **Timeseries (Metric)** Query:

- sum:siddes.api.request.count{\*}.as\_rate()

#### Widget 2: 5xx count (backend)

Type: **Timeseries (Logs)** Log query:

```
@event:api_request @status:[500 TO 599]
```

Compute: count

#### Widget 3: Error rate % (5xx / all)

Type: **Timeseries (Logs) + Formula** Query A (errors):

```
@event:api_request @status:[500 TO 599]
```

Query B (total):

```
@event:api_request
```

Formula:

- 100 \* A / B

#### Widget 4: Latency p95 (overall)

Type: **Timeseries (Metric)** Query:

- p95:siddes.api.request.latency\_ms{\*}

### Widget 5: Top failing routes (5xx)

Type: **Top List (Logs)** Query:

```
@event:api_request @status:[ 500 TO 599 ]
```

Group by: route Compute: count Top: 10

### Widget 6: Top slow routes (p95 latency)

Type: **Table (Metric)** Query:

- p95:siddes.api.request.latency\_ms{\*} by {route}

Sort by: p95 desc Limit: 15

## Dashboard 2 - Siddes: Endpoint Explorer

### Widget 1: Route table (RPS, p95)

Type: **Table (Metric)** Queries:

- RPS: sum:siddes.api.request.count{\*}.as\_rate() by {route}
- p95: p95:siddes.api.request.latency\_ms{\*} by {route}

Tip: In table settings, show both columns and sort by p95 or RPS depending on incident.

### Widget 2: Status breakdown for a selected route

Type: **Timeseries (Logs)** Query:

```
@event:api_request @route:$route
```

Group by: status Compute: count

Add a dashboard variable named \$route (type: facet route).

### Widget 3: Sample logs for selected route (with request\_id)

Type: **Log Stream** Query:

```
@event:api_request @route:$route
```

Columns to show:

- request\_id
- method
- status
- latency\_ms
- side

## Dashboard 3 - Siddes: Auth & Session

**Widget 1: 401/403 on auth endpoints**

Type: **Timeseries (Logs)** Query:

```
@event:api_request @path:/api/auth/* @status:(401 OR 403)
```

Compute: count Group by: status (optional)

**Widget 2: Write-method 401 outside auth (CSRF/cookie drift detector)**

Type: **Timeseries (Logs)** Query:

```
@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*
```

Compute: count

**Widget 3: Top routes causing write 401**

Type: **Top List (Logs)** Query:

```
@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*
```

Group by: route Compute: count Top: 10

**Dashboard 4 - Siddes: Media Pipeline****Widget 1: sign-upload success vs fail**

Type: **Timeseries (Logs) + Formula A (fail)**:

```
@event:api_request @path:"/api/media/sign-upload" @status:[400 TO 599]
```

B (total):

```
@event:api_request @path:"/api/media/sign-upload"
```

Formula:

- $100 * A / B$  (fail %)

**Widget 2: commit fail %**

Type: **Timeseries (Logs) + Formula A (fail)**:

```
@event:api_request @path:"/api/media/commit" @status:[400 TO 599]
```

B (total):

```
@event:api_request @path:"/api/media/commit"
```

Formula:

- $100 * A / B$

**Widget 3: commit p95 latency**

Type: **Timeseries (Metric)** Query:

- `p95:siddes.api.request.latency_ms{route:/api/media/commit}`

## Widget 4: Edge /m/\* status (Cloudflare)

Type: **Timeseries (Logs)** Query (assuming Logpush tags from the runbook):

```
service:siddes-edge @http.request.uri.path:/m/*
```

Group by: status (field name depends on your CF dataset; pick the status attribute you see) Compute: count

If your CF logs don't have service:siddes-edge, search ddsource:cloudflare and adjust.

## Dashboard 5 - Siddes: Inbox Reliability

### Widget 1: Inbox send error count

Type: **Timeseries (Logs)** Query:

```
@event:api_request @route:"/api/inbox/thread/:thread_id" @status:[400 TO 599]
```

Compute: count

### Widget 2: Inbox send p95 latency

Type: **Timeseries (Metric)** Query:

- p95:siddes.api.request.latency\_ms{route:/api/inbox/thread/:thread\_id}

### Widget 3: Inbox send failures by status

Type: **Top List (Logs)** Query:

```
@event:api_request @route:"/api/inbox/thread/:thread_id" @status:[400 TO 599]
```

Group by: status Compute: count Top: 10

## 3) Monitor Pack (Alerts)

Create these monitors in **Monitors** → **New Monitor**.

### Monitor 1 (PAGE): Backend 5xx spike

Type: **Log Monitor** Query:

```
logs("@event:api_request @status:[500 TO 599]").index("*").rollup("count").last("5m") > 20
```

Message (paste):

- “Backend 5xx spike. Check Golden Signals → Top failing routes. Pull sample request\_id and search logs.”

### Monitor 2 (PAGE): Core media commit failures

Type: **Log Monitor** Query:

```
logs("@event:api_request @path:\"/api/media/commit\" @status:[400 TO 599]").index("*").rollup("count").last("10m") > 10
```

Message:

- “Media commit failing. Check Media Pipeline dashboard. Likely R2/Worker/token-gate or upstream error.”

### **Monitor 3 (NOTIFY): Write 401 spike outside auth**

Type: **Log Monitor** Query:

```
logs( "@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE)
      -@path:/api/auth/*").index("*").rollup("count").last("10m") > 25
```

Message:

- “Write 401 spike outside auth. This often indicates cookie/CSRF/origin drift. Check Auth & Session dashboard.”

### **Monitor 4 (PAGE): Core route latency regression (p95)**

Type: **Metric Monitor** Query:

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 1500
```

Evaluation: last 10 minutes Message:

- “Feed p95 latency high. Check Endpoint Explorer sorted by p95. Pull slow request\_ids.”

If your feed route differs, update the `route:` tag to match your `route` facet value.

### **Monitor 5 (NOTIFY): Rate limit spike (429)**

Type: **Log Monitor** Query:

```
logs( "@event:api_request @status:429").index("*").rollup("count").last("10m") > 50
```

Message:

- “429 spike. Check which routes are being throttled; tune limits if needed.”

### **Monitor 6 (PAGE): Edge /m/\* 5xx (Cloudflare)**

Type: **Log Monitor** Query (adjust status attribute name if needed):

```
logs("service:siddes-edge @http.request.uri.path:/m/* @http.response.status_code:[500 TO
599]").index("*").rollup("count").last("10m") > 20
```

Message:

- “Edge /m 5xx spike. Check Media Pipeline dashboard and Cloudflare health.”

## **4) “How to use this” in a real incident (60 seconds)**

1) Open **Golden Signals** 2) Click the top failing `route` 3) In logs, copy a `request_id` 4) Search `request_id:<id>` and read:

- `status, latency_ms, route, side`

5) Fix starts from **where** it fails:

- auth 401 writes → CSRF/cookie/origin config

- media commit fails → R2/Worker token gating / upstream
- edge /m 403 → token mismatch or media privacy gate

If you want a “one dashboard only” version (super minimal), use Dashboard 1 only + Monitors 1-3.

# Siddes Observability Fire Drill (30 minutes, beginner-proof)

File: docs/OBSERVABILITY\_FIRE\_DRILL.md (source: sd\_602\_observability\_fire\_drill\_and\_next\_pack\_apply\_helper\_20260125\_054811.sh)

Goal: prove you can go from **a user complaint** → **a request\_id** → **the exact failing route/status/latency** in Datadog in under 2 minutes.

This drill is safe. It does **not** require code changes.

## 0) Pick targets (no guessing)

This drill needs two URLs:

- **WEB\_BASE** (your Vercel app URL) - used to test Next logs (inbox stub + /api/health)
- **API\_BASE** (your Django backend URL) - used to test backend api\_request logs

If you don't know them, they are already in your go-live docs:

- docs/GO\_LIVE\_DO\_VERCEL\_CLOUDFLARE\_RUNBOOK.md (look for SD\_INTERNAL\_API\_BASE=)

## 1) Prove request\_id correlation (Backend)

Run (replace the URL if needed):

```
curl -s -D - "$API_BASE/api/auth/csrf" -o /dev/null | sed -n 's/^X-Request-ID: //p'
```

Copy the request id.

In Datadog → Logs Explorer:

- Search: @event:api\_request @request\_id:<PASTE\_ID>
  - You should see exactly one log line.
  - Verify fields exist: path, status, latency\_ms, method
- Done when you can search by request\_id and find the backend log.

## 2) Prove route templating works (cardinality control)

Run (random id is fine; it can 404/401 - we only care about the route template):

```
curl -s -o /dev/null -w "%{http_code}"
```

```
" \"$API_BASE/api/postaaaaaaaaaaaa/like\"
```

In Datadog logs:

- Search: @event:api\_request @path:/api/post/\*/like
- Confirm your pipeline sets route=/api/post/:post\_id/like

■ Done when route is templated (not a million unique paths).

### 3) Prove the “CSRF/cookie drift detector” (write 401 outside auth)

Send an unauthenticated write request to a known write endpoint:

```
curl -s -o /dev/null -w "%{http_code}"
" -X POST \"$API_BASE/api/postaaaaaaaaaaaa/like\" -H "content-type: application/json" -d
'{}'
```

Expected: **401** (write guard)

In Datadog logs:

- Search: @event:api\_request @status:401 @method:POST -@path:/api/auth/\*
- You should see the request.

■ Done when the query above finds your test request.

### 4) Media pipeline sanity (backend)

Test sign-upload without auth:

```
curl -s -o /dev/null -w "%{http_code}"
" -X POST \"$API_BASE/api/media/sign-upload\" -H "content-type: application/json" -d '{}'
```

Expected: **401** or **403** (depending on your auth rules)

In Datadog logs:

- Search: @event:api\_request @path:"/api/media/sign-upload"

■ Done when you can see the request and status.

### 5) Prove Next (Vercel) logs exist for inbox stub

Run:

```
curl -s -o /dev/null -w "%{http_code}\n" "$WEB_BASE/api/inbox/threads"
```

In Datadog logs:

- Search for the string: /api/inbox/threads
- Confirm you can find the log entry from Vercel drain.

■ Done when Datadog shows Vercel logs for inbox.

## 6) Test monitors without breaking prod

You cannot safely “force 5xx” in production. Instead, for each monitor:

- Click **Test Notifications** (Datadog UI)
- Confirm you receive the notification/pager

Optional safe test:

- Temporarily lower a threshold (for 5 minutes) so a single test request trips it, then revert.
- Done when you can receive a test notification and you know where to find the route/request\_id evidence.

## 7) You are now production-debuggable

If a user reports “X didn’t work,” you can: 1) get request\_id (from header or response JSON) 2) search in logs 3) see route/status/latency 4) jump to dashboards for the wider picture

# Siddes Observability - Vercel/Next Pack (Inbox Stub + /api/health)

File: docs/OBSERVABILITY\_VERCEL\_NEXT\_PACK.md (source: sd\_602\_observability\_fire\_drill\_and\_next\_pack\_apply\_helper\_20260125\_054811.sh)

Important repo fact:

- docs/INBOX\_BACKEND\_CONTRACT.md says inbox is currently a **backend stub** using **Next.js route handlers**:
- frontend/src/app/api/inbox/threads/route.ts
- frontend/src/app/api/inbox/thread/[id]/route.ts

That means:

- Django @event:api\_request logs will **NOT** include inbox traffic.
- Inbox observability must come from **Vercel drained logs**.

This guide makes inbox debuggable in Datadog **without code changes**.

## 1) Find a Vercel inbox log entry (proof)

In Datadog Logs Explorer, search for:

```
/api/inbox/threads
```

Open one result and look at the attributes panel. Identify these 4 fields (names vary depending on your drain format):

- request id (often requestId or request\_id)
- path/url (often path or url)
- status code (often statusCode or status)
- duration in ms (often durationMs or duration\_ms)

## 2) Create facets (once)

For each attribute you found above:

- Click the attribute → **Create facet**

Create facets for:

- next\_request\_id (whatever the request id attribute is)
- next\_path
- next\_status

- next\_duration\_ms

Tip: If the log already uses path/status/duration, reuse those; do not duplicate.

### 3) Create 2 Next log-based metrics

#### A) Next requests count

In Logs Explorer query:

```
/api/inbox/
```

Generate metric:

- Name: siddes.next.request.count
- Type: count
- Tags: include your path/route field and status field (whatever you have)

#### B) Next duration distribution

Same query:

```
/api/inbox/
```

Generate metric:

- Name: siddes.next.request.duration\_ms
- Type: distribution
- Measure attribute: your duration-ms field

### 4) Build “Inbox (Next)” dashboard

#### Widget 1: Inbox RPS

Metric timeseries:

- sum:siddes.next.request.count{\*}.as\_rate()

#### Widget 2: Inbox errors

Log timeseries:

- query: /api/inbox/ @statusCode: [500 TO 599] (adjust to your status field)
- compute: count

#### Widget 3: Inbox p95 latency

Metric timeseries:

- p95:siddes.next.request.duration\_ms{\*}

## Widget 4: Top failing inbox routes

Top list (logs):

- query: /api/inbox/ @statusCode: [ 500 TO 599 ] (adjust)
- group by: path field
- compute: count

## 5) Add 2 Inbox monitors (Next)

### Monitor A (PAGE): Inbox 5xx spike (Next)

Log monitor:

- query: /api/inbox/ filtered to 5xx (adjust field)
- threshold: > 10 in 5 minutes

### Monitor B (NOTIFY): Inbox p95 high (Next)

Metric monitor:

- p95(siddes.next.request.duration\_ms{\*}) > 1500 for 10 minutes

## 6) Correlation workflow (when inbox fails)

When the UI says “send failed”:

- Go to Datadog logs, search /api/inbox/thread/
- Use the Vercel request id field (often requestId) to find the exact invocation
- If you also see a backend request\_id for other calls in the same session, you can correlate by time window

This closes the inbox observability gap without touching Django.

# Siddes Support Code Workflow (No Code Changes)

File: docs/OBSERVABILITY\_SUPPORT\_CODE\_WORKFLOW.md (source: sd\_603\_support\_code\_workflow\_docs\_apply\_helper\_20260125\_055031.sh)

This workflow turns “it didn’t work” into **one searchable key** in Datadog within 60 seconds, without collecting user content.

Siddes already provides the key for backend calls:

- X-Request-ID response header (and backend logs include `request_id`)
- Next proxy often injects `requestId` into JSON responses

## 1) What counts as a “Support Code”

### A) API Support Code = `request_id`

- The `request_id` is the fastest path to the exact server log line.
- It’s safe to share because it does **not** include user content.

### B) UI Support Code = “error digest”

- Next error boundaries can show `error.digest` (if/when exposed in UI).
- For now, use `request_id` for almost everything.

## 2) The 60-second operator loop (the whole point)

1) Get a `request_id` (see Section 3) 2) In Datadog logs search:

- `@event:api_request @request_id:<PASTE>`

3) Read:

- `route (or path), status, latency_ms, side`

4) Decide quickly:

- 401/403 → auth/session/CSRF/config drift
- 429 → rate limiting
- 5xx → backend/downstream issue
- 2xx but user says it failed → client/UI/state issue (capture via screenshots + timing; client telemetry later)

## 3) How to get `request_id` (Beginner-proof)

### Option A (Best): Web browser DevTools

1) Open Siddes in Chrome / Edge / Safari (desktop) 2) Open **DevTools** → **Network** 3) Reproduce the failure (click Like / Send / Upload) 4) Click the failing request in Network list 5) Find **X-Request-ID** in **Response Headers** 6) Copy it → that is the Support Code

If the response is JSON, also look for a `requestId` field in the response body.

### Option B (iPhone/iPad): Safari Web Inspector (works, but a bit setup-y)

1) On iPhone: **Settings** → **Safari** → **Advanced** → **Web Inspector** → **ON** 2) Connect iPhone to a Mac via cable 3) On Mac: Safari → **Settings** → **Advanced** → “**Show Develop menu**” → **ON** 4) On Mac Safari menu: **Develop** → <Your iPhone> → <Your Siddes tab> 5) Use Network panel → copy **X-Request-ID**

### Option C (No devtools): Use Datadog by time window

If you cannot extract `request_id`, you can still locate it:

1) Ask for:

- exact time (to the minute) + timezone
- what action they did (like / send / upload / login)
- which Side they were in (public/friends/close/work)

2) In Datadog logs:

- filter to `@event:api_request @status:[400 TO 599]`
- narrow by `@route:` (or `@path:`) based on the action (examples below)
- narrow time window ( $\pm 2$  minutes)

3) Pick the matching log line(s) and use `request_id` from there going forward.

#### Action → likely route/path filters

- Like: `@path:/api/post/*/like` or `@route:/api/post/:post_id/like`
- Post create: `@path:/api/post` (or your create endpoint)
- Media sign: `@path: "/api/media/sign-upload"`
- Media commit: `@path: "/api/media/commit"`
- Inbox send (Next stub): search Vercel logs for `/api/inbox/thread/` (see `docs/OBSERVABILITY_VERCEL_NEXT_PACK.md`)
- Auth failures: `@path:/api/auth/*`

## 4) What to do after you find the `request_id`

### If status is 401/403

- Open **Auth & Session** dashboard
- Check “Write-method 401 outside `/api/auth/*`” (cookie/CSRF/origin drift detector)

- Check /api/auth/csrf and /api/auth/me synthetic status (if enabled)

#### If status is 5xx

- Open **Golden Signals** → Top failing routes
- Confirm if it's one endpoint or many
- Pull 3 request\_ids to see if it's consistent (same route/status pattern)

#### If media is failing

- Open **Media Pipeline** dashboard
- If sign-upload ok but commit missing → likely browser PUT failed (invisible to backend)
- treat as client/network issue; collect time + action details and check edge logs

## 5) What NOT to ask users for (privacy)

Never request:

- screenshots of private messages/posts
- passwords/OTPs/tokens
- “send me your signed upload URL”
- email/phone unless strictly needed for account recovery (not debugging)

For debugging, you only need:

- time, action, side, and support code (request\_id) if possible.

# Siddes Support Response Macros (Copy/Paste)

File: docs/SUPPORT\_RESPONSE\_MACROS.md (source:  
sd\_603\_support\_code\_workflow\_docs\_apply\_helper\_20260125\_055031.sh)

Use these messages to gather the minimum info without collecting private content.

## A) First response (ask for Support Code + timing)

"Thanks - I can check the server logs. Please send: 1) The exact time it happened (and your timezone) 2) What you clicked (Like / Send / Upload / Login) 3) Which Side you were in (Public / Friends / Close / Work) 4) If you can: the **Support Code / Request ID** (it's in the X-Request-ID header on the failing request)."

## B) If they don't know how to get Request ID

"No worries - if you can't grab the Request ID, send the time (to the minute), what you clicked, and which Side you were in. I'll locate it in logs."

## C) Confirm you found it (reassuring + specific)

"Got it. I found your request in logs:

- Route: <route>
- Status: <status>
- Request ID: <request\_id>

We're tracking this now."

## D) Known classes (choose one)

### Auth/Session (401/403)

"This looks like an auth/session issue (401/403). We're checking cookie/CSRF configuration and will push a fix."

### Server error (5xx)

"This is a server-side error (5xx). We're investigating the failing endpoint and deploying a fix."

### Rate limit (429)

"This was rate-limited (429). We're tuning limits so normal usage doesn't get blocked."

### Media upload

"This looks like an upload pipeline issue. We're checking the sign/commit steps and edge delivery."

## **E) Closeout (with lightweight verification)**

“We deployed a fix. Please retry the same action. If it still fails, send the new Request ID and the time.”

# Siddes Incident Note Template (Copy/Paste)

File: docs/INCIDENT\_TEMPLATE.md (source:  
sd\_603\_support\_code\_workflow\_docs\_apply\_helper\_20260125\_055031.sh)

## Summary

- Start time:
- End time:
- Severity:
- Impact (what users saw):
- Scope (routes / services):

## Detection

- Alert/monitor name:
- First signal observed:
- Dashboard(s) used:

## Evidence (no PII)

- Example request\_ids:
- Top routes affected:
- Status patterns:
- Latency patterns:

## Root cause

- What broke:
- Why it broke:
- What changed (deploy/config):

## Fix

- What we did:
- When deployed:

## Verification

- What checks passed (healthz/readyz/auth/me/media):

## Prevention

- New monitor / threshold changes:

- Runbook updates:
- Follow-ups:

# Siddes Oncall Home (Single Pane)

File: docs/ONCALL\_HOME.md (source: sd\_604\_oncall\_home\_runbook\_apply\_helper\_20260125\_055726.sh)

This page is the “start here” for incidents. It assumes you have Datadog dashboards + monitors set up.

## 1) The one thing you ask for: Request ID

If a user reports “it didn’t work,” your first move is:

- get **Request ID** (X-Request-ID / requestId)

Then search:

- Datadog Logs → @event:api\_request @request\_id:<PASTE>

Docs:

- docs/OBSERVABILITY\_SUPPORT\_CODE\_WORKFLOW.md

## 2) Open these dashboards (in this order)

### A) Golden Signals (Prod)

Purpose: “Are we on fire?”

- RPS
- 5xx rate
- p95 latency
- top failing routes
- top slow routes

### B) Endpoint Explorer

Purpose: “Which route is broken / slow?”

- route table (rps + p95)
- status breakdown
- sample logs with request\_id

### C) Auth & Session

Purpose: “Did cookies/CSRF break?”

- 401/403 on auth endpoints
- write 401 outside auth detector

## D) Media Pipeline

Purpose: "Can users upload/view media?"

- sign-upload fail %
- commit fail %
- edge /m status breakdown

## E) Inbox Reliability

Purpose: "Are messages sending?"

- send fail count
- send p95 latency
- status breakdown

## 3) When an alert fires (what to do immediately)

### Alert: Backend 5xx spike

1) Golden Signals → Top failing routes 2) Click failing route → Endpoint Explorer 3) Pull 3 request\_ids from logs 4) Confirm if it's one route or many

### Alert: Write 401 spike outside auth

1) Auth & Session dashboard 2) Check recent deploy/config change (origins/cookies/CSRF) 3) Verify /api/auth/csrf + /api/auth/me health

### Alert: Media commit failing

1) Media Pipeline dashboard 2) Check sign-upload vs commit 3) Check edge /m/\* status 4) If sign ok but commit missing: treat as client PUT/network (collect time/action; investigate edge + client signals)

### Alert: Inbox 5xx (Next)

1) Inbox dashboard (Next) 2) Check Vercel logs for /api/inbox/ 3) Confirm whether any backend calls are also failing

## 4) How to write an incident note (no PII)

Use:

- docs/INCIDENT\_TEMPLATE.md

## 5) Fire drill (keep your skills sharp)

Run anytime:

- ./scripts/obs/fire\_drill.sh

**Docs:**

- `docs/OBSERVABILITY_FIRE_DRILL.md`

# Siddes Oncall Quickstart (Beginner)

File: docs/ONCALL\_QUICKSTART.md (source: sd\_604\_oncall\_home\_runbook\_apply\_helper\_20260125\_055726.sh)

If you only memorize one flow, memorize this:

## 1) User report → Request ID → Datadog search

1) Get X-Request-ID from the failing request (best) 2) Search Datadog logs: @event:api\_request @request\_id:<id> 3) Read status, route, latency\_ms

## 2) Fast classification

- 401/403 → Auth/session/CSRF/cookie drift
- 429 → Rate limiting
- 5xx → Server/downstream
- 2xx but UI says it failed → Client/UI state mismatch

## 3) Where to click next

- 5xx spike → Golden Signals → Endpoint Explorer → sample request\_ids
- auth drift → Auth & Session dashboard
- media → Media Pipeline dashboard + edge /m/\*
- inbox → Inbox dashboard (Next logs)

## 4) Don't collect private content

Never ask users for:

- passwords, OTPs, tokens
- post/DM text
- signed URLs

Only ask for:

- time (to the minute) + timezone
- action they took
- which Side
- request id (support code)

# Siddes Observability Master Checklist (Rebuild From Scratch)

File: docs/OBSERVABILITY\_MASTER\_CHECKLIST.md (source:  
sd\_605\_observability\_master\_checklist\_apply\_helper\_20260125\_060017.sh)

This checklist is designed for a beginner. Follow it top-to-bottom. It assumes **no code changes**.

## Phase 0 - Ground truth (already in repo)

■ Backend emits JSON request logs for `/api/*` with `event="api_request"` and `request_id` ■  
Backend returns `X-Request-ID` and exposes it to browsers via CORS ■ Next proxy forwards  
`x-request-id` and often injects `requestId` into JSON ■ Go-live docs define API base and health  
checks

Definition of Done:

- You can see `@event:api_request` in your log tool (Datadog).

## Phase 1 - Centralize logs (single pane)

### 1.1 DigitalOcean App Platform → Datadog Logs

Do:

- Enable DO log forwarding to Datadog

DoD:

- Datadog Logs shows `@event:api_request`

### 1.2 Vercel → Datadog Logs (log drain)

Do:

- Enable Vercel log drains to Datadog

DoD:

- Datadog Logs includes Vercel runtime logs
- Searching `/api/inbox/threads` returns at least one entry after a request

### 1.3 Cloudflare Logpush (/m/\*) → Datadog Logs

Do:

- Enable Cloudflare Logpush filtered to paths starting `/m/`

DoD:

- Datadog Logs shows edge logs for /m/\*
- You can see status breakdown (200/302/401/403/5xx)

## Phase 2 - Privacy controls (non-negotiable)

### 2.1 Hash or drop viewer identifiers

Do:

- Use Sensitive Data Scanner (or pipeline) to hash viewer
- Drop dev\_viewer

DoD:

- New logs never show viewer=me\_123 (raw)
- viewer is hashed or removed

### 2.2 Redact or drop query

Do:

- Redact attribute query to [redacted\_query] or drop it

DoD:

- New logs never include full query strings

### 2.3 Route templating (cardinality control)

Do:

- Create pipelines that set route templates for dynamic paths:
  - /api/post/:post\_id/like
  - /api/post/:post\_id/reply
  - /api/post/:post\_id/replies
  - /api/post/:post\_id/quote
  - /api/inbox/thread/:thread\_id
  - /m/:key (edge side)

DoD:

- Grouping by route produces a small, stable set of values (not thousands)

## Phase 3 - Metrics (log-based)

Create 2 metrics from logs: 1) siddes.api.request.count (count) 2) siddes.api.request.latency\_ms (distribution on latency\_ms)

DoD:

- You can graph RPS from siddes.api.request.count
- You can graph p95 latency from siddes.api.request.latency\_ms

## Phase 4 - Dashboards

Build these dashboards: 1) Golden Signals 2) Endpoint Explorer 3) Auth & Session 4) Media Pipeline 5) Inbox Reliability

DoD:

- In Golden Signals, you can identify top failing route in < 30 seconds
- In Endpoint Explorer, you can drill down and copy a request\_id

## Phase 5 - Alerts (Monitors)

Create these monitors:

- PAGE: backend 5xx spike (count threshold)
- PAGE: media commit failures
- NOTIFY: write-method 401 spike outside /api/auth/\*
- PAGE: feed p95 latency > 1500ms
- NOTIFY: 429 spike
- PAGE: edge /m/\* 5xx spike

DoD:

- “Test notification” works for each monitor
- Each alert message tells you which dashboard to open and what to look for

## Phase 6 - Fire Drill (prove it works)

Run:

- ./scripts/obs/fire\_drill.sh

DoD:

- You can find each test request in Datadog with the suggested queries

- You can locate a request\_id and read status/latency/route

## Phase 7 - Support workflow (real-world usage)

Use:

- docs/OBSERVABILITY\_SUPPORT\_CODE\_WORKFLOW.md
- docs/SUPPORT\_RESPONSE\_MACROS.md

DoD:

- For a user report, you can respond with:
- “I found your request in logs: route/status/request\_id”

within 2 minutes

## Phase 8 - Oncall start page

Pin:

- docs/ONCALL\_HOME.md
- docs/ONCALL\_QUICKSTART.md

DoD:

- In an incident, you know exactly where to click first.

If you completed Phases 1-8, Siddes is production-debuggable without violating privacy.

# Observability (Production Debugging, Privacy-First)

File: docs/OBSERVABILITY.md (source: sd\_614\_observability\_md\_rewrite\_helper\_20260125\_075509.sh)

This is the **main entry point** for Siddes observability.

If you are a beginner, start here:

- Open `docs/OBS_PACK_INDEX.md` and follow the read order.
- Pin `docs/ONCALL_HOME.md` for real incidents.

## What Siddes already gives you (baseline)

### Backend (Django)

- Every `/api/*` response includes `X-Request-ID`.
- Requests to `/api/*` emit JSON-line logs via `siddes.api` logger:  
`event, request_id, viewer, method, path, status, latency_ms, side`
- Write guard exists in production to block unauthenticated writes (default allowlist is `/api/auth/*`).

### Frontend (Next.js)

- Core routes include `error.tsx` boundaries to show calm failure + retry:
- Feed, Post detail, Sets, Invites, Inbox
- The `/api` proxy generates/forwards `x-request-id` and often injects `requestId` into JSON responses.

## The 5-minute questions (the whole purpose)

In 5 minutes, oncall must be able to answer:

- Are 5xx errors spiking? Which endpoint(s)? Which service (Vercel vs Django vs Cloudflare edge)?
- Is p95 latency spiking? Which endpoint(s)?
- Are auth failures (401/403) spiking on write endpoints (cookie/CSRF/origin drift)?
- Are media failures happening (sign-upload/commit) and/or `/m/*` failing at the edge?
- Is inbox failing (note: inbox is currently a Next stub, so it's Vercel logs, not Django logs)?

## Privacy-first logging rules (non-negotiable)

### Required stored fields (canonical)

Store only request metadata needed for debugging:

- `request_id, route (templated), method, status, latency_ms, side`

- plus operational tags: service, env, release, client\_version (if available)

### **NEVER store (PII / secrets / content)**

- passwords, OTPs, session cookies, auth/refresh tokens, Authorization
- email/phone/username/name
- post bodies, comments, inbox/DM text, bios, search queries
- signed media URLs, media keys, raw IP addresses

### **Mandatory ingestion transforms (no code changes)**

1) **Hash or drop** viewer (do not index/store raw me\_123) 2) **Drop/redact** query 3) Convert dynamic paths → **templated** route (cardinality control), e.g.:

- /api/post/:post\_id/like
- /api/inbox/thread/:thread\_id
- /m/:key

Policy docs:

- docs/OBSERVABILITY\_RETENTION\_ACCESS\_POLICY.md
- docs/OBSERVABILITY\_PRIVACY\_AUDIT\_CHECKLIST.md

### **Event taxonomy (minimal)**

Server/edge:

- api.request (Django request logs; derived from event="api\_request")
- next.request (Vercel runtime logs for Next route handlers - required for Inbox)
- media.edge\_request (Cloudflare /m/\* logs)

Client (optional, log-only lane):

- client.error.boundary
- client.net.request\_failed
- client.media.put\_failed

### **Datadog setup + dashboards + alerts (copy/paste)**

#### **ClickOps implementation**

- docs/OBSERVABILITY\_CLICKOPS\_DATADOG.md

#### **Dashboards & monitors pack**

- docs/OBSERVABILITY\_DATADOG\_DASHBOARDS\_ALERTS.md

### SLOs + burn-rate alerts (smarter paging)

- docs/OBSERVABILITY\_SLOS\_V0.md
- docs/OBSERVABILITY\_BURN\_RATE\_ALERTS.md

## How to debug a user report in 60 seconds (Support Code workflow)

**Always ask for Request ID** (X-Request-ID / requestId). Then search:

- Datadog Logs: @event:api\_request @request\_id:<PASTE>

Use:

- docs/OBSERVABILITY\_SUPPORT\_CODE\_WORKFLOW.md
- docs/SUPPORT\_RESPONSE\_MACROS.md

Incident notes:

- docs/INCIDENT\_TEMPLATE.md

## Readiness gates + drills (make sure it works)

### Before every prod deploy (do not deploy blind)

- docs/OBSERVABILITY\_LAUNCH\_GATE.md
- scripts/go\_live\_observability\_gate.sh

### Fire drill (prove request\_id + route templating)

- docs/OBSERVABILITY\_FIRE\_DRILL.md
- scripts/obs/fire\_drill.sh

### Incident drills (practice real scenarios)

- docs/OBSERVABILITY INCIDENT\_DRILLS.md
- scripts/obs/incident\_drills.sh

### “What am I seeing?” and “What do I fix first?”

- docs/OBSERVABILITY\_KNOWN\_FAILURE\_SIGNATURES.md
- docs/OBSERVABILITY\_FIX\_PLAYBOOK\_MAP.md

## **Rebuild everything anytime**

- `docs/OBSERVABILITY_MASTER_CHECKLIST.md`

If you only pin ONE doc for incidents:

- `docs/ONCALL_HOME.md`

# Siddes Burn-Rate Alerts (Smarter paging, less noise)

File: docs/OBSERVABILITY\_BURN\_RATE\_ALERTS.md (source:  
sd\_608\_burn\_rate\_alerts\_apply\_helper\_20260125\_061745.sh)

Burn-rate alerts page you when you're **spending error budget too fast**. This avoids constant paging on tiny spikes.

This v0 uses:

- Log monitors (5xx) as the “bad events”
- Metric monitors (p95 latency) as “user pain”
- Route templates (`route`) from your pipelines

## 1) Concepts (simple)

If your SLO is 99.5% over 7 days, your error budget is 0.5%. Burn rate means: “how fast am I consuming that 0.5% budget?”

We'll implement a practical approximation:

- **Fast burn**: big outage now → page quickly
- **Slow burn**: steady degradation → notify and fix before budget is gone

## 2) Availability burn alerts (5xx)

We use **two monitors per core route group**:

### A) FAST BURN (PAGE)

Trigger when 5xx is high in a short window. Example (Feed):

- Window: 5 minutes
- Threshold: 5xx count > 20

Datadog Log Monitor query:

```
logs "@event:api_request @route:/api/feed @status:[500 TO 599]".rollup("count").last("5m") > 20
```

### B) SLOW BURN (NOTIFY)

Trigger when 5xx is elevated over a longer window. Example (Feed):

- Window: 60 minutes
- Threshold: 5xx count > 60

## Datadog Log Monitor query:

```
logs("@event:api_request @route:/api/feed @status:[500 TO 599]").rollup("count").last("1h") > 60
```

These are starter thresholds. After 2 weeks of traffic data, tune them by baseline.

## 3) Latency burn alerts (p95)

Latency doesn't consume "error budget", but users feel it like outages. Use the same pattern:

### A) FAST LATENCY (PAGE)

Example (Feed):

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 2000
```

Evaluate: last 10 minutes

### B) SLOW LATENCY (NOTIFY)

Example (Feed):

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 1500
```

Evaluate: last 60 minutes

## 4) Core route groups (copy/paste templates)

### Feed

- route: /api/feed

FAST 5xx:

```
logs("@event:api_request @route:/api/feed @status:[500 TO 599]").rollup("count").last("5m") > 20
```

SLOW 5xx:

```
logs("@event:api_request @route:/api/feed @status:[500 TO 599]").rollup("count").last("1h") > 60
```

FAST p95:

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 2000
```

SLOW p95:

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 1500
```

### Auth: csrf + me

If you have route tags for these, use them. Otherwise use path. CSRF 5xx:

```
logs("@event:api_request @path:'/api/auth/csrf' @status:[500 TO 599]").rollup("count").last("5m") > 10
```

**ME 5xx:**

```
logs("@event:api_request @path:'/api/auth/me' @status:[500 TO 599]").rollup("count").last("5m") > 10
```

**Latency:**

```
p95(siddes.api.request.latency_ms{route:/api/auth/me}) > 1000
```

### **Media: commit (critical)**

FAST commit failures (PAGE):

```
logs("@event:api_request @path:'/api/media/commit' @status:[400 TO 599]").rollup("count").last("10m") > 10
```

SLOW commit failures (NOTIFY):

```
logs("@event:api_request @path:'/api/media/commit' @status:[400 TO 599]").rollup("count").last("1h") > 30
```

**Latency:**

```
p95(siddes.api.request.latency_ms{route:/api/media/commit}) > 2500
```

### **Auth drift detector (write 401 spike outside auth) - NOTIFY**

```
logs("@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*").rollup("count").last("10m") > 25
```

This is not burn-rate, but it's your best "configuration broke" early warning.

## **5) Inbox (Next stub) burn alerts (Vercel logs)**

Inbox is currently Next route handlers in your repo, so use Next metrics:

- siddes.next.request.count
- siddes.next.request.duration\_ms

FAST 5xx (PAGE) - adjust fields:

- Log monitor filtered to 5xx on /api/inbox/ for 5m

FAST p95 (PAGE):

```
p95(siddes.next.request.duration_ms{*}) > 2000
```

SLOW p95 (NOTIFY):

```
p95(siddes.next.request.duration_ms{*}) > 1500
```

## **6) Alert messages (always include these)**

Every monitor should say: 1) Which dashboard to open first (Golden Signals / Endpoint Explorer / Media / Auth) 2) What query to run 3) "Pull 3 request\_ids and search them"

Example message: “Feed errors burning budget. Open Golden Signals → Endpoint Explorer. Search logs: @event:api\_request @route:/api/feed @status:[ 500 TO 599 ] Pull 3 request\_ids and inspect status/latency.”

After 2 weeks of traffic, we can switch from raw counts to true burn-rate % monitors.

# Siddes - Fix Playbook Map (Signature → Evidence → Repo Knobs → Verify)

File: docs/OBSERVABILITY\_FIX\_PLAYBOOK\_MAP.md (source: sd\_613\_fix\_playbook\_map\_apply\_helper\_20260125\_064746.sh)

This document links the most common production failure signatures to the **exact Siddes docs and config knobs** to check. It assumes **no application code changes** - only ClickOps/config and verification scripts.

If you have a `request_id`, always start here:

- Datadog Logs: @event:api\_request @request\_id:<PASTE>

## 0) Universal verification tools (use these first)

### A) Proxy / health sanity (post deploy)

Run:

```
./scripts/post_deploy_smoke.sh <APP_ORIGIN> <API_ORIGIN> [MEDIA_TEST_URL]
```

Source: scripts/post\_deploy\_smoke.sh (tests /healthz, /readyz, /api/health, /api/auth/me, /api/auth/csrf, feed, optional media)

### B) Observability Gate (before deploy)

Run:

```
WEB_BASE=<APP_ORIGIN> API_BASE=<API_ORIGIN> ./scripts/go_live_observability_gate.sh
```

Source: docs/OBSERVABILITY\_LAUNCH\_GATE.md

### C) Observability Fire Drill (prove request\_id + route templating)

Run:

```
WEB_BASE=<APP_ORIGIN> API_BASE=<API_ORIGIN> ./scripts/obs/fire_drill.sh
```

Source: docs/OBSERVABILITY\_FIRE\_DRILL.md

## 1) Signature: backend\_not\_configured / proxy\_fetch\_failed / no backend logs

### What users see

- Login page shows `proxy_fetch_failed`
- API responses show `{ ok:false, error:"backend_not_configured" }`
- Inbox might work (Next stub), but real backend calls fail

### Datadog evidence

- Vercel logs show errors for `/api/*`

- Django logs have **no matching** api\_request for those calls

### Where in Siddes docs

- docs/FRONTEND\_PROXY\_DIAG.md
- docs/GO\_LIVE\_DO\_VERCEL\_CLOUDFLARE\_RUNBOOK.md
- docs/GO\_LIVE\_CLICKOPS\_DO\_VERCEL\_CLOUDFLARE.md
- Template: ops/deploy/vercel.env.example

### Config knobs (what to check)

On Vercel:

- SD\_INTERNAL\_API\_BASE=https://api.yourdomain.com

On backend (DigitalOcean):

- backend must be reachable and /healthz returns 200

### Fix steps (ClickOps)

1) Vercel → Project → Settings → Environment Variables:

- set SD\_INTERNAL\_API\_BASE exactly to your backend origin

2) Confirm backend health directly:

- curl -i https://<backend>/healthz

3) Confirm proxy health:

- GET https://<app>/api/\_diag should show ok:true and healthz.ok:true

### Verify

Run:

```
./scripts/post_deploy_smoke.sh https://<app> https://<api>
```

## 2) Signature: Write actions return 401/403 (CSRF/cookie/origin drift)

### What users see

- “Can’t post”, “can’t like”, “send fails”
- Works in some browsers, not others
- Often begins right after a domain/proxy/CDN change

### Datadog evidence

Write-method 401 spike outside auth:

```
@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*
```

### Auth 401/403:

```
@event:api_request @path:/api/auth/* @status:(401 OR 403)
```

### Where in Siddes docs

- docs/PROD\_DOMAINS\_COOKIES\_CSRF.md
- docs/GO\_LIVE\_DO\_VERCEL\_CLOUDFLARE\_RUNBOOK.md
- docs/GO\_LIVE\_CLICKOPS\_DO\_VERCEL\_CLOUDFLARE.md
- Template: ops/deploy/backend.env.prod.example

### Config knobs (what to check)

#### On DigitalOcean (Django env vars):

- DJANGO\_ALLOWED\_HOSTS=api.yourdomain.com
- DJANGO\_CSRF\_TRUSTED=https://app.yourdomain.com \*(must include Vercel/custom app origins)\*
- Optional: SIDDES\_COOKIE\_DOMAIN=.yourdomain.com \*(only if you want cross-subdomain cookies)\*
- DJANGO\_DEBUG=0

#### On Vercel:

- SD\_INTERNAL\_API\_BASE=https://api.yourdomain.com

### Fix steps (ClickOps)

1) Update DO env vars (allowed hosts + csrf trusted) 2) Redeploy backend 3) Confirm:

- /api/auth/csrf returns {ok:true}
- /api/auth/me returns valid JSON {ok:true, ...} (even if viewer null when logged out)

4) If you use Cloudflare in front of your app domain:

- ensure it is not stripping cookies/headers

### Verify

Run:

```
./scripts/post_deploy_smoke.sh https://<app> https://<api>
```

### 3) Signature: Media upload stuck (R2 CORS / PUT blocked) - sign-upload ok, commit missing

#### What users see

- Upload spinner never finishes

- Avatar upload fails
- Browser console shows CORS error (PUT blocked)

### Datadog evidence

- You see sign-upload requests:
- @event:api\_request @path: "/api/media/sign-upload"
- But commit is missing (or commit drops sharply):
- @event:api\_request @path: "/api/media/commit"
- Often: sign-upload OK, then nothing (because browser PUT never reaches your servers)

### Where in Siddes docs

- docs/GO\_LIVE\_DO\_VERCEL\_CLOUDFLARE\_RUNBOOK.md (R2 bucket + CORS)
- docs/GO\_LIVE\_CLICKOPS\_DO\_VERCEL\_CLOUDFLARE.md (R2 ClickOps)
- docs/MEDIA\_R2.md (architecture)
- (Reference) overlay history: sd\_503\_cloudflare\_r2\_cors\_autopilot\_\* (optional helper)

### Config knobs (what to check)

#### On Cloudflare R2 bucket CORS:

- Methods: PUT, GET, HEAD
- Origins: your **exact** Vercel/custom app origins (no trailing slash)
- Headers: content-type

### Fix steps (ClickOps)

1) Cloudflare Dashboard → R2 → Bucket → CORS 2) Set allowed origins exactly to your app origin(s) 3)  
Save

### Verify

- Perform an upload in the app
- In Datadog:
- sign-upload happens
- commit happens
- Optional: run smoke + do a manual upload test

### 4) Signature: /m/\* returns 401/403 for allowed users (Worker token secret mismatch)

### What users see

- Post loads but image is broken

- Some viewers see media, others see 403
- Incognito gets blocked (expected), but logged-in viewer also blocked (bad)

### Datadog evidence

- Backend sign/commit calls succeed
- Edge logs show /m/\* returning 401/403 spikes
- Backend logs may show GET /api/media/url?key=... activity (if used)

### Where in Siddes docs

- docs/CLOUDFLARE\_MEDIA\_WORKER\_TOKEN\_GATE.md
- ops/deploy/cloudflare\_worker\_setup.md
- docs/MEDIA\_R2.md

### Config knobs (what to check)

#### On DigitalOcean (Django env vars):

- SIDDES\_MEDIA\_TOKEN\_SECRET=<secret>
- Optional: SIDDES\_MEDIA\_PRIVATE\_TTL=600

#### On Cloudflare Worker secrets:

- MEDIA\_TOKEN\_SECRET must match SIDDES\_MEDIA\_TOKEN\_SECRET

### Fix steps (ClickOps)

1) Confirm the DO secret value 2) In Worker project:

- npx wrangler secret put MEDIA\_TOKEN\_SECRET
- deploy Worker

3) Confirm Cloudflare route exists for /m/\*

### Verify

Follow docs/CLOUDFLARE\_MEDIA\_WORKER\_TOKEN\_GATE.md "What to test"

- Private-side media loads for allowed viewer
- Incognito fails (or fails after TTL)

### 5) Signature: /m/\* returns 404 / bypasses worker / edge logs missing

#### What users see

- All media links 404

- Or media loads only in dev, not in prod

### Datadog evidence

- No Cloudflare /m/\* logs OR status is 404/5xx at edge
- Backend logs might show redirects in dev fallback, but production path not routed

### Where in Siddes docs

- docs/CLOUDFLARE\_MEDIA\_WORKER\_TOKEN\_GATE.md (Route section)
- ops/deploy/cloudflare\_worker\_setup.md
- docs/DOMAINS\_AND\_CLOUDFLARE\_DNS\_PLAN.md

### Config knobs (what to check)

#### In Cloudflare:

- Worker route: https://yourdomain.com/m/\* (or https://app.yourdomain.com/m/\* depending on your plan)

### Fix steps (ClickOps)

1) Cloudflare → Workers & Pages → Routes 2) Add route for /m/\* to the Worker 3) Ensure DNS for app domain points correctly

### Verify

- Hit a real media URL:
- curl -I https://app.yourdomain.com/m/<key>
- expect 200/302 (depending on caching/token behavior)

## 6) Signature: 429 spikes (rate limiting too tight)

### What users see

- “Try again” / intermittent failures
- More frequent for heavy usage or bursts

### Datadog evidence

```
@event:api_request @status:429
```

Group by route to see where.

### Where in Siddes docs

- docs/THROTTLING.md (scopes + env overrides)
- docs/TIME\_LIMIT\_PLAYBOOK.md (if timeouts also present)

## Config knobs (what to check)

On **DigitalOcean (Django env vars)**, tune:

- SIDDES\_THROTTLE\_INBOX\_THREADS
- SIDDES\_THROTTLE\_INBOX\_THREAD
- SIDDES\_THROTTLE\_INBOX\_SEND
- and other scope env vars listed in `docs/THROTTLING.md`

## Fix steps (ClickOps)

- 1) Identify affected route(s) in Datadog
- 2) Find scope in `docs/THROTTLING.md`
- 3) Increase that scope env var
- 4) Redeploy backend

## Verify

- 429 rate drops
- normal traffic no longer throttled

## 7) Signature: 5xx spike / outages

### What users see

- “Server error”
- Feed doesn’t load

### Datadog evidence

```
@event:api_request @status:[500 TO 599]
```

Golden Signals dashboard shows spike, Endpoint Explorer shows top failing routes.

### Where in Siddes docs

- `docs/GO_LIVE_MASTER_RUNBOOK.md`
- `docs/DEPLOYMENT_GATES.md`
- `docs/GO_LIVE_ISSUE_REGISTER.md`

## Fix steps (first moves)

- 1) Is it one route or many? 2) If many routes:
  - suspect DB/infrastructure saturation
  - restart/scale backend if needed
- 3) If one route:
  - rollback the last deploy/config change

## Verify

Run:

```
./scripts/post_deploy_smoke.sh https://<app> https://<api>
```

## 8) Signature: Inbox broken (Next/Vercel stub)

Repo fact:

- Inbox is currently served by Next route handlers:
- frontend/src/app/api/inbox/threads/route.ts
- frontend/src/app/api/inbox/thread/[id]/route.ts

### What users see

- Inbox list doesn't load
- Send fails quickly

### Datadog evidence

- Vercel logs show /api/inbox/\* errors
- Django @event:api\_request logs may show nothing for inbox

### Where in Siddes docs

- docs/INBOX\_BACKEND\_CONTRACT.md
- docs/INBOX\_STUB\_SMOKE\_DEMO.md
- docs/INBOX\_VISIBILITY\_STUB.md
- docs/OBSERVABILITY\_VERCEL\_NEXT\_PACK.md (your observability pack)

### Fix steps (ClickOps)

1) Confirm Vercel log drain is active (so you can see inbox logs) 2) Check Vercel deployment + env vars 3)  
Use the Next inbox dashboard/monitors from the Next Pack doc

## Verify

Run:

```
curl -i https://<app>/api/inbox/threads
```

and confirm the request appears in Datadog.

## 9) Signature: “We’re blind” (logs missing)

### What you see

- No @event:api\_request logs
- Dashboards empty

### Where in Siddes docs

- docs/OBSERVABILITY.md
- docs/OBSERVABILITY\_CLICKOPS\_DATADOG.md
- docs/OBSERVABILITY\_LAUNCH\_GATE.md
- docs/OBSERVABILITY\_MASTER\_CHECKLIST.md

### Fix steps (ClickOps)

1) Verify DO log forwarding 2) Verify Vercel drain 3) Verify Datadog index/pipelines still match @event:api\_request 4) Re-run:

```
WEB_BASE=https://<app> API_BASE=https://<api> ./scripts/go_live_observability_gate.sh
```

If you want a final “everything in one page” version: use docs/ONCALL\_HOME.md as your pin.

# Siddes Observability Incident Drill Pack (Beginner, 45 minutes)

File: docs/OBSERVABILITY INCIDENT DRILLS.md (source: sd\_611\_observability\_incident\_drill\_pack\_apply\_helper\_20260125\_062916.sh)

These drills make you confident you can handle real incidents. They are **safe**: they generate normal test traffic and expected error codes (401/404), and teach you where to click + what to search.

You do **not** need code changes.

Prereq:

- Logs are flowing to Datadog (@event:api\_request works)
- Dashboards exist (Golden Signals, Endpoint Explorer, Auth, Media, Inbox)
- Fire drill already passed (./scripts/obs/fire\_drill.sh)

## Setup (2 minutes)

You need:

- WEB\_BASE = your Vercel URL (e.g. https://siddes.app)
- API\_BASE = your Django backend URL (e.g. https://api.siddes.app)

If you don't know them, check your go-live runbook docs (look for SD\_INTERNAL\_API\_BASE=).

## Drill 1 - Auth Drift Detector (Write 401 spike outside auth)

**Goal:** prove you can detect “posting/liking broke due to cookie/CSRF/origin drift” quickly.

### Step A: Generate the signal (safe)

Run an unauthenticated write request (expected 401 in prod):

```
curl -s -o /dev/null -w "%{http_code}"
" -X POST "$API_BASE/api/post/aaaaaaaaaaaaaaaa/like" -H "content-type: application/json" -d
'{}'
```

### Step B: Verify the signal in Datadog (exact query)

In Datadog Logs Explorer search:

```
@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*
```

### Step C: Where to click

- Dashboard: **Auth & Session**
- Look at “Write-method 401 outside auth”

- Look at “Top routes causing write 401”

#### **Step D: What you would do in a real incident**

- Confirm whether /api/auth/csrf and /api/auth/me are healthy
- If write 401 spikes suddenly after deploy → suspect:
  - cookie domain mismatch
  - CSRF trusted origins mismatch
  - proxy/Cloudflare header/cookie stripping

#### **Pass condition**

- You can find the request in logs and see it in the Auth dashboard.

### **Drill 2 - Media Pipeline Break (sign/commit failing)**

**Goal:** prove you can see sign/commit requests and interpret failures without looking at user content.

#### **Step A: Generate sign-upload signal (safe)**

Unauth sign-upload (expected 401/403):

```
curl -s -o /dev/null -w "%{http_code}"
" -X POST "$API_BASE/api/media/sign-upload" -H "content-type: application/json" -d '{}'
```

#### **Step B: Generate commit signal (safe)**

Unauth commit (expected 401/403/400 depending on rules):

```
curl -s -o /dev/null -w "%{http_code}"
" -X POST "$API_BASE/api/media/commit" -H "content-type: application/json" -d '{}'
```

#### **Step C: Verify in Datadog**

Search:

```
@event:api_request @path:"/api/media/sign-upload"
```

and:

```
@event:api_request @path:"/api/media/commit"
```

#### **Step D: Where to click**

- Dashboard: **Media Pipeline**
- sign-upload fail %
- commit fail %
- commit p95 latency
- edge /m/\* status breakdown (if CF logs enabled)

### How to reason (real incident)

- sign-upload fails → auth/config or backend issue
- sign-upload OK but commit missing → likely browser PUT to R2 failed (client/network), which backend can't see
- commit fails → auth rules, validation, upstream/R2 edge issues

### Pass condition

- You can find both sign and commit requests in logs and see them on Media dashboard.

## Drill 3 - Inbox Break (Next/Vercel stub observability)

**Goal:** prove you can debug inbox even if it never hits Django (because it's Next route handlers).

### Step A: Generate inbox requests (safe)

```
curl -s -o /dev/null -w "%{http_code}"
" ${WEB_BASE}/api/inbox/threads"
curl -s -o /dev/null -w "%{http_code}"
" ${WEB_BASE}/api/inbox/thread/test"
```

### Step B: Verify in Datadog (Vercel logs)

Search:

- /api/inbox/threads
- /api/inbox/thread/

If you created Next metrics per `docs/OBSERVABILITY_VERCEL_NEXT_PACK.md`, also verify:

- `siddes.next.request.count` exists
- `p95(siddes.next.request.duration_ms{*})` graphs

### Step C: Where to click

- Inbox Reliability dashboard (Next)
- Vercel logs stream filtered to `/api/inbox/`

### Pass condition

- You can find inbox logs in Datadog and see request durations/statuses.

## After each drill: Practice the “support response”

Use:

- `docs/SUPPORT_RESPONSE_MACROS.md`

Write a pretend reply to a user using the request\_id approach:

- “Found your request in logs: route/status/request\_id.”

### Bonus: Practice writing an incident note (5 minutes)

Use:

- docs/INCIDENT\_TEMPLATE.md

Fill it with:

- start/end time (approx)
- what monitor would have fired
- 3 example request\_ids from the drill

### If any drill fails (what to fix)

- No backend logs → DO log forwarding / Datadog index/pipeline
- No Vercel inbox logs → Vercel drain
- No edge /m/\* logs → Cloudflare Logpush filter/tagging
- route missing → route templating pipeline rules not applied
- privacy regression → Sensitive Data Scanner rules not matching

# Siddes Known Failure Signatures (Datadog “What am I seeing?” Guide)

File: docs/OBSERVABILITY\_KNOWN\_FAILURE\_SIGNATURES.md (source: sd\_612\_known\_failure\_signatures\_apply\_helper\_20260125\_064045.sh)

This is a beginner-friendly map from **symptom** → **Datadog evidence** → **first fix to try**. No code changes required.

Use this during incidents.

## 0) Always start with `request_id` (if you have it)

Datadog search:

- `@event:api_request @request_id:<id>`

Read:

- `route/path, status, latency_ms, side`

If you don't have `request_id`:

- filter by time window ( $\pm 2$  minutes) + route + status.

## 1) CSRF / Cookie / Origin Drift (classic “everything broke after deploy”)

### User symptom

- “I can’t like / post / send”
- “It keeps saying not authenticated”
- Works for some browsers, not others

### Datadog signatures

A) Write-method 401 spike outside auth:

```
@event:api_request @status:401 @method:(POST OR PUT OR PATCH OR DELETE) -@path:/api/auth/*
```

B) Auth endpoints behaving oddly:

```
@event:api_request @path:/api/auth/* @status:(401 OR 403)
```

### Dashboard to open

- **Auth & Session**

### First fix to try (config)

- Verify cookie domain + CSRF trusted origins match current domains
- Verify Cloudflare isn't stripping cookies/headers
- Verify /api/auth/csrf returns 200 and sets cookies correctly
- Re-run scripts/go\_live\_observability\_gate.sh

## 2) Backend outage / downstream failure (true 5xx)

### User symptom

- Feed doesn't load
- Any action returns "server error"

### Datadog signatures

A) Global 5xx spike:

```
@event:api_request @status:[ 500 TO 599 ]
```

B) Find top failing routes (group by route)

### Dashboards

- Golden Signals → Endpoint Explorer

### First fix to try

- Check if it's one endpoint or many
- If many routes slow + 5xx: suspect DB / infra
- If one route: suspect a specific regression; roll back if needed

## 3) Media uploads broken (R2 / Worker / CORS / token gate)

### User symptom

- "Upload stuck"
- "Image doesn't appear"
- "Avatar upload fails"
- Media loads sometimes, sometimes 403

### Datadog signatures

A) sign-upload fails:

```
@event:api_request @path:"/api/media/sign-upload" @status:[ 400 TO 599 ]
```

B) commit fails:

```
@event:api_request @path:"/api/media/commit" @status:[ 400 TO 599 ]
```

### C) Edge /m failures (Cloudflare): Search logs:

- service:siddes-edge @http.request.uri.path:/m/\*  
@http.response.status\_code:(401 OR 403 OR 5\*)

(adjust field names)

### Dashboard

- Media Pipeline

### Interpretation

- sign-upload fails → auth/config/back-end
- sign-upload OK but commit missing → browser PUT to R2 likely failed (client/network)
- commit fails with 4xx → validation/auth mismatch
- /m 403 spikes → token gate mismatch or side/privacy enforcement issue
- /m 5xx spikes → worker/R2 outage

### First fix to try (config)

- Re-check R2 CORS allowed origins match **exactly** your web origin
- Confirm Worker env vars and token gate config
- Confirm media commit route has correct auth requirements

## 4) Inbox broken (Next/Vercel stub issues)

Repo fact:

- Inbox is currently served by Next route handlers (backend\_stub). Django logs won't show it.

### User symptom

- Inbox list doesn't load
- Send fails immediately

### Datadog signatures

A) Vercel logs show /api/inbox/\* errors

- Search: /api/inbox/

B) If you created metrics:

- siddes.next.request.count drop or spikes
- p95(siddes.next.request.duration\_ms{\*}) spike

## Dashboard

- Inbox Reliability (Next)

### First fix to try

- Check Vercel deployment changes
- Check Vercel environment variables (API base / any feature flags)
- If inbox is expected to call backend, confirm proxy routes and CORS

## 5) Rate limiting (429)

### User symptom

- “Sometimes it works, sometimes it says try again”
- Heavy users hit failures

### Datadog signatures

```
@event:api_request @status:429
```

Group by route to see where it's happening.

## Dashboard

- Endpoint Explorer

### First fix to try

- Increase limits for the affected route(s)
- Add targeted throttling rather than global

## 6) Performance regression (no errors, but slow)

### User symptom

- App “spins” but eventually loads
- Works late at night but slow at peak

### Datadog signatures

- p95 latency rising across multiple routes
- Top slow routes table shows broad slowdown

### Dashboards

- Golden Signals (p95) → Endpoint Explorer (sort by p95)

### First fix to try

- Look for deploy marker correlation
- If only one route: optimize/rollback that route
- If many routes: check DB connections/saturation

## 7) “It succeeded but user says it didn’t” (client state mismatch)

### User symptom

- “Like didn’t change”
- “Message sent but didn’t appear”

### Datadog signature

- Request returns 2xx quickly in logs, but UI didn’t update

### What to do

- Ask for Request ID + time
- Confirm server returned 2xx
- Then treat as client cache/state issue:
- collect screenshots (without private content)
- check for client-side errors (once Part 2 client lane is enabled)

If you want the next item: a “Fix Playbook Map” that links each signature to the exact go-live docs and config knobs in your repo.

# Siddes Launch Readiness - Observability Gate (Run Before Every Prod Deploy)

File: docs/OBSERVABILITY\_LAUNCH\_GATE.md (source:  
sd\_610\_observability\_launch\_gate\_apply\_helper\_20260125\_062727.sh)

This gate prevents “we deployed and now we’re blind” incidents. It requires **no code changes**.

Run this before each prod deploy and after any:

- Vercel drain changes
- Cloudflare logpush/worker changes
- backend logging/middleware changes
- Datadog pipeline/index changes

## 1) Required checks (pass/fail)

### A) Backend logs present

**Pass if:**

- Datadog has recent @event:api\_request logs in the last 10 minutes.

### B) Request ID correlation works

**Pass if:**

- curl \$API\_BASE/api/auth/csrf returns an X-Request-ID
- searching Datadog by that request\_id returns the log line

### C) Privacy controls still active

**Pass if:**

- viewer is hashed or missing (not me\_123)
- query is redacted or missing
- no auth secrets appear in logs (Authorization/Cookie/Bearer)

### D) Route templating still active

**Pass if:**

- grouping by route shows a small stable set
- /api/post/:post\_id/like templating still works

### E) Vercel inbox logs present (Next stub)

**Pass if:**

- a request to \$WEB\_BASE/api/inbox/threads shows up in Datadog logs

#### F) Monitors are armed

Pass if:

- monitors exist for:
- backend 5xx spike
- media commit failures
- write 401 drift detector
- feed p95 latency
- edge /m 5xx (if CF logs enabled)

#### 2) What to do if any check fails

1) Stop the deploy (or deploy only to staging) 2) Fix the blind spot:

- log forwarding / drain / logpush
- pipeline rules (hash/redact/route)

3) Re-run the gate 4) Only then deploy

#### 3) Automated helper script

Use:

- scripts/go\_live\_observability\_gate.sh

It runs safe curl checks and prints the exact Datadog queries to verify.

# Siddes Observability Privacy Audit Checklist (Monthly)

File: docs/OBSERVABILITY\_PRIVACY\_AUDIT\_CHECKLIST.md (source:  
sd\_609\_observability\_privacy\_audit\_checklist\_apply\_helper\_20260125\_062008.sh)

This checklist ensures observability stays **privacy-safe** over time, especially as code and infrastructure change.

Do this once per month (or before a major launch).

## A) Verify sensitive fields are not stored

### A1) Viewer identifier is hashed or removed

In Datadog Logs Explorer:

- Query: @event:api\_request
- Open a recent log line

Check:

- viewer is NOT present, OR it is hashed (not me\_123)
- dev\_viewer is NOT present

■ Pass criteria:

- No raw viewer IDs appear in any logs.

### A2) Query strings are redacted/dropped

Check:

- query is either missing or equals [redacted\_query]

■ Pass criteria:

- No user-entered query/search strings appear.

### A3) No auth secrets appear

Search these patterns (each should return **zero** results):

- Authorization:
- Bearer
- Set-Cookie:
- csrftoken=
- sessionid=
- refresh

- token

■ Pass criteria:

- Zero hits in logs.

## B) Verify cardinality controls (prevents accidental data leaks)

### B1) Route templating is active

In logs, facet/group by route.

Check:

- You see a small stable set like:
- /api/post/:post\_id/like
- /api/inbox/thread/:thread\_id
- /api/media/commit
- etc.

■ Pass criteria:

- You do NOT see thousands of unique values.

### B2) Raw path is not used for dashboards

Check your dashboards:

- They group by route, not raw path.

■ Pass criteria:

- Dashboards are stable and not exploding in unique paths.

## C) Verify retention settings (privacy + cost)

### C1) Raw logs retention

In Datadog:

- Logs → Configuration → Indexes (or Retention)

Check:

- Raw logs retention is set to  $\leq 7$  days (recommended)
- Errors/sampled logs are retained appropriately, but not forever

**■ Pass criteria:**

- Raw logs do not retain beyond your policy.

**C2) Telemetry DB retention (if enabled)**

In backend env:

- SIDDES\_TELEMETRY\_RETENTION\_DAYS is set (default 30)
- Purge job exists and ran recently

**■ Pass criteria:**

- Telemetry rows older than retention are gone.

**D) Verify access controls (least privilege)****D1) Who can see raw logs?**

Confirm:

- Support team does **not** have raw log access.
- Oncall/Engineers have access only as needed.

**■ Pass criteria:**

- Access aligns with “tiers” policy.

**D2) Sensitive Data Scanner still enabled**

Confirm:

- Scanner group still matches @event:api\_request
- Hash rule for viewer is active
- Redact rule for query is active

**■ Pass criteria:**

- Rules are active and applied to new logs.

**E) Verify the incident workflow stays privacy-safe****E1) Support Code workflow uses Request IDs only**

Check:

- Support macros ask for time/action/side/request\_id
- They do not ask for content or tokens

**■ Pass criteria:**

- No user content is requested.

**E2) Incident template contains no PII**

Check:

- Example request\_ids only
- No user messages/post text

**■ Pass criteria:**

- Incident notes are clean.

**F) Regression checks after deploy changes**

Whenever you change:

- logging pipelines
- Cloudflare worker settings
- Vercel drains
- backend middleware

Run:

- `./scripts/obs/fire_drill.sh`

**■ Pass criteria:**

- You can still find logs by request\_id and route templating still works.

If any check fails: 1) Disable indexing for the sensitive field immediately (stop the bleed) 2) Fix the pipeline rules (hash/redact) 3) Reduce retention temporarily 4) Re-run the fire drill

# Siddes Observability Retention & Access Policy (v0)

File: docs/OBSERVABILITY\_RETENTION\_ACCESS\_POLICY.md (source:  
sd\_609\_observability\_privacy\_audit\_checklist\_apply\_helper\_20260125\_062008.sh)

This is a short, enforceable policy you can share with anyone who asks “what do we store?”

## 1) Data retained

### A) Operational logs (Datadog)

- Purpose: production debugging
- Stored fields: request metadata only (route, status, latency, request\_id)
- Identity: viewer is hashed or removed
- Content: never stored

Retention:

- Raw logs: **7 days** max (recommended)
- Aggregated metrics: 30-90 days

### B) Telemetry events (DB)

- Purpose: counts-only product signals
- Stored fields: event name + viewer id + timestamp
- Retention: **30 days** (default)

### C) Incident notes

- Purpose: learning + prevention
- Must not contain PII or user content
- Retention: as long as needed (safe because no PII)

## 2) Access tiers

### Tier 1 - Oncall/SRE

- Can search raw logs by request\_id
- Can view route/status/latency
- Cannot access secrets like HMAC salt

### Tier 2 - Engineers

- Dashboards by default

- Raw logs time-bound when needed for incidents

### **Tier 3 - Support**

- No raw logs
- Uses Support Code workflow (request\_id/time/action/side)

### **3) Prohibited collection**

Never collect or store in observability systems:

- tokens, cookies, Authorization headers
- usernames/emails/phones
- post bodies, comments, inbox messages
- signed media URLs or media keys
- raw IP addresses

### **4) Enforcement**

- Sensitive Data Scanner rules must remain enabled
- Monthly privacy audit checklist must be completed

See:

- [docs/OBSERVABILITY\\_PRIVACY\\_AUDIT\\_CHECKLIST.md](#)

# Siddes Observability SLOs v0 (No Code Changes)

File: docs/OBSERVABILITY\_SLOS\_V0.md (source: sd\_607\_slos\_v0\_apply\_helper\_20260125\_061239.sh)

SLOs make alerts smarter and calmer. This v0 uses only what Siddes already logs:

- `@event:api_request with status and latency_ms`
- route templating (`route`) from your Datadog pipelines

This is **not** “enterprise SRE.” It’s a simple starter so you stop paging on noise.

## 1) Definitions (simple)

### Availability SLI

For an endpoint group:

- **Good** = status is 2xx or 3xx
- **Bad** = status is 5xx
- (4xx usually means user/auth/validation; not “service down”)

Availability SLI formula:

- `good / total where total = good + bad`

### Latency SLI

For an endpoint group:

- We track **p95 latency** and keep it below a threshold.

## 2) SLOs (starter targets)

These are intentionally forgiving for an early-stage product. Tighten later.

### A) Feed SLO (core)

- Scope: `route:/api/feed`
- Availability: **99.5%** over 7 days (5xx only)
- Latency: p95 < **1500ms** over 30 minutes

### B) Auth SLO (core)

- Scope: `/api/auth/me, /api/auth/csrf`
- Availability: **99.7%** over 7 days (5xx only)
- Latency: p95 < **800ms** over 30 minutes

### C) Media SLO (core)

- Scope: /api/media/sign-upload, /api/media/commit
- Availability: **99.0%** over 7 days (5xx only)
- Latency: p95 < **2000ms** over 30 minutes
- Extra guardrail: /api/media/commit **4xx spike** monitor (can indicate auth/cors problems)

### D) Inbox SLO (core user pain)

Because inbox is currently a Next stub in your repo, split it:

- **Inbox (Next):** /api/inbox/\* from Vercel logs/metrics
- **Inbox (Backend):** only if/when those endpoints hit Django

Targets:

- Availability: **99.0%** over 7 days (5xx only)
- Latency: p95 < **1500ms** over 30 minutes

## 3) Error budget (what it means in plain English)

Example: 99.5% availability over 7 days:

- Total minutes in 7 days =  $7 * 24 * 60 = 10080$
- 0.5% budget = 50.4 minutes of “5xx time” allowed (across that endpoint group)

This is **not permission to be down** - it's a tool to decide when to:

- stop shipping features
- focus on stability/perf

## 4) How to implement in Datadog (minimal)

### 4.1 Create a “5xx rate” monitor per SLO scope

For each scope, you can monitor:

- 5xx count
- or 5xx rate % (better once traffic is steady)

Example (Feed 5xx spike) - Log Monitor:

```
logs("@event:api_request @route:/api/feed @status:[500 TO 599]").rollup("count").last("5m") >
10
```

### Example (Feed latency regression) - Metric Monitor:

```
p95(siddes.api.request.latency_ms{route:/api/feed}) > 1500
```

### 4.2 Convert to SLOs later (optional)

Datadog has a dedicated SLO product that uses monitors as the SLI source. Once traffic is steady:

- define SLOs from these monitors
- track error budget burn

### 5) Tightening plan (what to do later)

After 2 weeks of data: 1) Look at typical p95 for each core route 2) Set latency threshold to ~2x typical p95  
3) Lower availability error budgets gradually (99.0 → 99.5 → 99.7)

### 6) What not to do

- Don't page on 4xx globally (it's noisy)
- Don't page on p95 for non-core endpoints
- Don't group by raw path (use `route` templates only)

If you want, the next step is “burn-rate alerts” (fast + slow burn) for each SLO, still without code changes.

# Siddes Observability Pack Index (Read Order)

File: docs/OBS\_PACK\_INDEX.md (source: sd\_606\_obs\_pack\_index\_apply\_helper\_20260125\_060359.sh)

If you are a beginner, read these in order.

## 1) What we are building (spec)

1) docs/OBSERVABILITY\_SPEC\_V1.md

- The rules: what we log, what we never log, what we must answer in 5 minutes

## 2) Set it up in the real world (ClickOps)

2) docs/OBSERVABILITY\_CLICKOPS\_DATADOG.md

- Connect DO → Datadog logs
- Connect Vercel → Datadog logs
- Connect Cloudflare /m/\* → Datadog logs
- Hash/redact sensitive fields
- Route templating

## 3) Build dashboards + alerts (copy/paste)

3) docs/OBSERVABILITY\_DATADOG\_DASHBOARDS\_ALERTS.md

- Create 2 log-based metrics
- Build dashboards (Golden Signals, Endpoint Explorer, Auth, Media, Inbox)
- Add monitors (5xx spike, media, auth drift, latency)

## 4) Prove it works (fire drill)

4) docs/OBSERVABILITY\_FIRE\_DRILL.md 5) scripts/obs/fire\_drill.sh

## 5) Close the Next/Vercel inbox observability gap (important!)

6) docs/OBSERVABILITY\_VERCEL\_NEXT\_PACK.md

- Inbox is a Next route-handler stub in your repo
- You must use Vercel drained logs to debug inbox

## 6) Support workflow (real incidents)

- 7) docs/OBSERVABILITY\_SUPPORT\_CODE\_WORKFLOW.md
- 8) docs/SUPPORT\_RESPONSE\_MACROS.md
- 9) docs/INCIDENT\_TEMPLATE.md

## 7) Oncall start page (pin this)

- 10) docs/ONCALL\_HOME.md
- 11) docs/ONCALL\_QUICKSTART.md

## 8) The “rebuild from scratch” checklist

- 12) docs/OBSERVABILITY\_MASTER\_CHECKLIST.md

### Apply helpers (scripts you ran)

- sd\_600\_... → writes spec + clickops + playbook
- sd\_601\_... → writes Datadog dashboards/alerts doc
- sd\_602\_... → writes fire drill + Next inbox pack + helper script
- sd\_603\_... → writes support code workflow + macros + incident template
- sd\_604\_... → writes oncall home + quickstart
- sd\_605\_... → writes master checklist
- sd\_606\_... → writes this index

If you only pin ONE doc, pin:

- docs/ONCALL\_HOME.md

# Siddes Oncall Observability Playbook (Beginner-friendly)

File: docs/ONCALL\_OBSERVABILITY\_PLAYBOOK.md (source: sd\_600\_observability\_docs\_datadog\_apply\_helper\_20260125\_053723.sh)

This playbook assumes you have centralized logs + the Golden Signals dashboard.

## 0) The one thing you always ask for: Request ID

Siddes gives you `request_id / X-Request-ID` for `/api/*`.

When someone reports:

- “like didn’t work”
- “post failed”
- “upload stuck”

Your first move is: **get the Request ID** (from the response JSON `requestId` or the `X-Request-ID` header).

Then: 1) Search logs for that `request_id` 2) Read `route, status, latency_ms`

## 1) Triage flow (2 minutes)

### 1) Open **Golden Signals**

- 5xx spike? → jump to **Endpoint Explorer** and sort by errors
- p95 spike? → sort by slow endpoints

### 2) Check **Auth & Session**

- write 401 spike outside `/api/auth/*` often means CSRF/cookie/origin misconfig

### 3) Check **Media Pipeline**

- sign-upload ok but commit missing → likely browser PUT failed (client signal helps later)

### 4) Check **Edge /m/**

- If `/m/*` returns 401/403 unexpectedly → token gate or auth mismatch

## 2) Common incidents and what they look like

### A) “Users can’t post / everything is 401”

Symptoms:

- write-method 401 spike outside /api/auth/\*
- /api/auth/csrf failing or cookies not being set

What to do:

- confirm Vercel → API origin setup matches go-live runbook
- confirm Cloudflare proxy settings didn't change cookie/headers

### B) “Media uploads broken”

Symptoms:

- /api/media/sign-upload ok but /api/media/commit fails
- or commit success drops sharply
- edge /m/\* status spikes

What to do:

- check R2 CORS settings (origin match exact)
- check Worker token-gate behavior and env vars

### C) “Inbox send broken”

Symptoms:

- /api/inbox/thread/:thread\_id 5xx/timeout spike
- or Vercel runtime errors if inbox is Next-side

What to do:

- check whether errors are in Vercel logs or Django logs (topology split)

## 3) How to answer a user report in 60 seconds

If you have the request\_id:

- “I see your request hit the server at <time>, returned <status>, route <route>. We’re fixing <cause>.”

If you don’t have request\_id:

- “Please send the Support Code shown on the error screen.” (use request\_id/digest)

## 4) What “good” looks like (baseline)

- 5xx near zero
- p95 latency under ~1.5s on core routes

- media commit success stable
- auth 401 mostly limited to unauthenticated reads, not write actions